

1. SUBRINGS OF QUOTIENT RINGS

We may have overlooked the case where the ambient ring is a quotient ring. This use case is important to some applications and it is discussed the paper [“Using SAGBI bases to compute invariants”](#) by Stillman and Tsai. If you aren’t already aware, this paper’s authors are the package’s original programmers. In the paper, they specifically mention a specific implementation which is an earlier version of the package.

The code is actually close to being capable of handling subrings of quotient rings. After looking into it, I think it should work correctly in all cases where a finite Sagbi basis happens to exist. However, a problem arises when a finite Sagbi basis doesn’t exist. In this case, it is possible for the algorithm to falsely succeed. The calculation they give in Example 1 exhibits this behavior:

```
i1 : gndR = kk[x,y, MonomialOrder => Lex];
```

```
i2 : I = ideal(x^2 - x*y);
```

```
o2 : Ideal of gndR
```

```
i3 : Q = gndR/I;
```

```
i4 : subR = sagbi subring {x};
```

```
i5 : gens subR
```

```
o5 = | x |
```

```
o5 : Matrix Q   $\begin{matrix} 1 & 1 \\ & \end{matrix}$  <--- Q
```

In this example, the algorithm terminated successfully because it has no way of knowing that $x^2 = xy$ inside of the ambient ring. Technically, one could argue that this isn’t a bug because it’s actually doing what it is supposed to do according to the specification. Notice that the Sagbi algorithm’s proof of correctness (Proposition 3) doesn’t apply when a finite Sagbi basis doesn’t exist. Fixing this “issue” is really a question of whether or not it is possible to improve upon the underlying algorithm.

In cases like the above example, it would be an improvement if there was a way to at least distinguish between failure and success. In section 2.2, Stillman and Tsai give some information about how this can be done. However, I find this part difficult to understand and I haven’t been able to apply successfully apply their solution. Another problem is that it complicates the design because there would be really be two types of failure, for a total of three possible outcomes:

- (1) Successfully found a finite Sagbi basis.
- (2) Failed to find a finite Sagbi basis before the limit.
- (3) (The new condition) Ran out of S-polynomial candidates.

I wonder if it’s possible to generate a useful partial Sagbi basis in case (3).